

# Boosting

Victor Kitov

[v.v.kitov@yandex.ru](mailto:v.v.kitov@yandex.ru)

# Linear ensembles

## Linear ensemble:

$$F_M(x) = f_0(x) - c_1 f_1(x) + \dots - c_M f_M(x)$$

**Regression:**  $\hat{y}(x) = F_M(x)$

**Binary classification:**  $\text{score}(y|x) = F_M(x)$ ,  $\hat{y}(x) = \text{sign } F_M(x)$

- Notation:  $f_1(x), \dots, f_M(x)$  are called *base learners*, *weak learners*, *base models*.
- Too expensive to optimize  $f_0(x), f_1(x), \dots, f_M(x)$  and  $c_1, \dots, c_M$  jointly for large  $M$ .
- Idea: optimize  $f_0(x)$  and then each pair  $(f_m(x), c_m)$  step-by-step.

# Forward stagewise additive modeling (FSAM)

**Input:**

- training dataset  $(x_n, y_n)$ ,  $n = 1, 2, \dots, N$
- loss function  $\mathcal{L}(f, y)$
- parametric form of base learner  $f_\theta(x)$
- the number of base learners  $M$ .

**Output:** approximation function  $F_M(x) = f_0(x) - \sum_{m=1}^M c_m f_m(x)$

# Forward stagewise additive modeling (FSAM)

- ❶ Fit initial approximation  $f_0(x) = \arg \min_f \sum_{n=1}^N \mathcal{L}(f(x_n), y_n)$
- ❷ For  $m = 1, 2, \dots, M$ :
  - find next best classifier


$$(c_m, f_m) := \arg \min_{f, c} \sum_{n=1}^N \mathcal{L}(F_{m-1}(x_n) - cf(x_n), y_n)$$

- reevaluate ensemble

$$F_m(x) := F_{m-1}(x) - c_m f_m(x)$$

# Dependency on $M$

Boosting overfits for high  $M$ :



## Comments

- $M$  should be determined by performance on validation set.
- Each step should be coarse to leave room for future base learners improvement:
  - initial approximation may be zero or constant
  - optimization can be coarse (just few steps)
  - base learner should be simple
    - such as trees of depth=1,2,3.
- For some loss functions (see Adaboost) we can solve minimization explicitly.
- For general loss functions gradient boosting should be used.

# Table of Contents

1 Adaboost

2 Gradient boosting

3 Extensions

# Adaboost (discrete version)

## Assumptions:

- binary classification task  $y \in \{+1, -1\}$
- $f_m(x) \in \{+1, -1\}$ , trainable in weighted dataset.
- classification is performed with
$$\hat{y} = \text{sign}\{f_0(x) + c_1 f_1(x) + \dots + c_M f_M(x)\}$$
- optimized loss is  $\mathcal{L}(F(x), y) = e^{-yF(x)}$

**Optimization in FSAM can be solved explicitly!**



# Adaboost (discrete version): algorithm

**Input:**

- training dataset  $(x_n, y_n)$ ,  $n = 1, 2, \dots, N$
- number of additive weak classifiers  $M$
- a family of weak classifiers  $f_m(x) \in \{+1, -1\}$ 
  - should be trainable on weighted datasets.

**Output:** composite classifier  $F_M(x) = \text{sign} \left( \sum_{m=1}^M c_m f_m(x) \right)$

## Adaboost (discrete version): algorithm

- ❶ Initialize observation weights  $w_n = 1/N$ ,  $n = 1, 2, \dots, N$ .
- ❷ for  $m = 1, 2, \dots, M$ :
  - ❶ fit  $f_m(x)$  to training data using weights  $w_n$
  - ❷ compute weighted misclassification rate:

$$E_m = \frac{\sum_{n=1}^N w_n \mathbb{I}[f_m(x_n) \neq y_n]}{\sum_{n=1}^N w_n}$$

- ❸ if  $E_m > 0.5$  or  $E_m = 0$ : terminate procedure.
- ❹ compute  $c_m = \frac{1}{2} \ln((1 - E_m)/E_m)$   $E_m < 0.5 \Rightarrow c_m > 0$
- ❺ increase all weights, where misclassification with  $f_m(x)$  was made:

$$w_n \leftarrow w_n e^{2c_m} = w_n \left( \frac{1 - E_m}{E_m} \right), \text{ for } n : f_m(x_n) \neq y_n$$

# Adaboost derivation

Set  $F_0(x) \equiv 0$ .

Apply FSAM for  $m = 1, 2, \dots, M$ :

$$\begin{aligned}(c_m, f_m) &= \arg \min_{c_m, f_m} \sum_{n=1}^N \mathcal{L}(F_{m-1}(x_n) + c_m f_m(x_n), y_n) \\&= \arg \min_{c_m, f_m} \sum_{n=1}^N e^{-y_n F_{m-1}(x_n)} e^{-c_m y_n f_m(x_n)} \\&= \arg \min_{c_m, f_m} \sum_{i=1}^N w_n^m e^{-c_m y_n f_m(x_n)}, \quad w_n^m := e^{-y_n F_{m-1}(x_n)}\end{aligned}$$

# Adaboost derivation

$$\begin{aligned}
 \sum_{n=1}^N w_n^m e^{-c_m y_n f_m(x_n)} &= \sum_{n: f_m(x_n)=y_n} w_n^m e^{-c_m} + \sum_{n: f_m(x_n) \neq y_n} w_n^m e^{c_m} \\
 &= e^{-c_m} \sum_{n: f_m(x_n)=y_n} w_n^m + e^{c_m} \sum_{n: f_m(x_n) \neq y_n} w_n^m \\
 &= e^{c_m} \sum_{n: f_m(x_n) \neq y_n} w_n^m + e^{-c_m} \sum_{n=1}^N w_n^m - e^{-c_m} \sum_{n: f_m(x_n) \neq y_n} w_n^m \\
 &= e^{-c_m} \sum_n w_n^m + (e^{c_m} - e^{-c_m}) \sum_{n: f_m(x_n) \neq y_n} w_n^m
 \end{aligned}$$

Since  $c_m \geq 0$ ,  $f_m(\cdot)$  should be found from

$$f_m(\cdot) = \arg \min_f \sum_{n=1}^N w_n^m \mathbb{I}[f(x_n) \neq y_n]$$

# Adaboost derivation

Denote  $G(c_m) = \sum_{n=1}^N w_n^m \exp(-c_m y_n f_m(x_n))$ . Then

$$\frac{\partial G(c_m)}{\partial c_m} = - \sum_{n=1}^N w_n^m e^{-c_m y_n f_m(x_n)} y_n f_m(x_n) = 0$$

$$- \sum_{n: f_m(x_n)=y_n} w_n^m e^{-c_m} + \sum_{n: f_m(x_n) \neq y_n} w_n^m e^{c_m} = 0$$

$$e^{2c_m} = \frac{\sum_{n: f_m(x_n)=y_n} w_n^m}{\sum_{n: f_m(x_n) \neq y_n} w_n^m}$$

$$c_m = \frac{1}{2} \ln \frac{\left( \sum_{n: f_m(x_n)=y_n} w_n^m \right) / \left( \sum_{n=1}^N w_n^m \right)}{\left( \sum_{n: f_m(x_n) \neq y_n} w_n^m \right) / \left( \sum_{n=1}^N w_n^m \right)} = \frac{1}{2} \ln \frac{1 - E_m}{E_m},$$

$$\text{where } E_m := \frac{\sum_{n=1}^N w_n^m \mathbb{I}[f_m(x_n) \neq y_n]}{\sum_{n=1}^N w_n^m}$$

## Adaboost derivation

Weights recalculation:

$$w_n^{m+1} \stackrel{\text{def}}{=} e^{-y_n F_m(x_n)} = e^{-y_n F_{m-1}(x_n)} e^{-y_n c_m f_m(x_n)}$$

Noting that  $-y_n f_m(x_n) = 2\mathbb{I}[f_m(x_n) \neq y_n] - 1$ , we can rewrite:

$$\begin{aligned} w_n^{m+1} &= e^{-y_n F_{m-1}(x_n)} e^{c_m(2\mathbb{I}[f_m(x_n) \neq y_n] - 1)} = \\ &= w_n^m e^{2c_m \mathbb{I}[f_m(x_n) \neq y_n]} e^{-c_m} \propto w_n^m e^{2c_m \mathbb{I}[f_m(x_n) \neq y_n]} \end{aligned}$$

Comments:

- We can remove common constants from weights.
- $w_n^{m+1} = w_n^m$  for correctly classified objects by  $f_m(x)$ .
- $w_n^{m+1} = w_n^m e^{2c_m}$  for incorrectly classified objects by  $f_m(x)$ .
  - so later classifiers will pay more attention to them

# Table of Contents

- 1 Adaboost
- 2 Gradient boosting
- 3 Extensions

# Motivation

- Problem: For general loss function  $L$  FSAM cannot be solved explicitly
- Analogy with function minimization: when we can't find optimum explicitly we use numerical methods
- Gradient boosting: numerical method for iterative loss minimization



# Gradient descent algorithm

$$L(w) \rightarrow \min_w, \quad g(w) = \nabla_w L(w), \quad w \in \mathbb{R}^N$$

Gradient descent:

```
initialize  $w$   
for  $m = 1, 2, \dots M$ :  
     $g(w) = \nabla_w L(w)$   
     $w = w - \varepsilon g(w)$ 
```

Gradient descent with modified step:

```
initialize  $w$   
for  $m = 1, 2, \dots M$ :  
     $g(w) = \nabla_w L(w)$   
     $c^* = \arg \min_{c > 0} L(w - cg(w))$   
     $w = w - c^* \Delta w$ 
```

## Gradient boosting intuition

$$L(F) = \sum_{n=1}^N \mathcal{L}(F^n) \rightarrow \min_F \quad F = [F^1, F^2, \dots, F^N]$$

$$\text{Gradient descent:} \quad F := F - c \nabla L(F)$$

$$\text{Pointwise gradient descent:} \quad F^n := F^n - c \nabla L(F) = F^n - c \nabla \mathcal{L}(F^n)$$

We want generalization to new  $x$ , so need functional approximation:

$$F(x) := F(x) - cf(x)$$

$$f(x_n) \approx \nabla \mathcal{L}(F(x_n)) \quad n = 1, 2, \dots, N$$

# Gradient boosting

- Now consider

$$L(f(x_1), \dots, f(x_N)) = \sum_{n=1}^N \mathcal{L}(f(x_n), y_n) \rightarrow \min_{f(\cdot)}$$

- Gradient descent performs pointwise optimization, but we need generalization, so we optimize in space of functions.
- Gradient boosting = modified gradient descent in function space:
  - find gradients:  $g(x_n) = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \big|_{r=f^{m-1}(x_n)}$
  - fit base learner  $f_m(x)$  to  $\{(x_n, g(x_n))\}_{n=1}^N$

# Gradient boosting

**Input:** training dataset  $(x_n, y_n)$ ,  $n = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$  and the number  $M$  of successive additive approximations.

- 1 Fit initial approximation  $f_0(x)$  (might be taken  $f_0(x) \equiv 0$ )

# Gradient boosting

**Input:** training dataset  $(x_n, y_n)$ ,  $n = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$  and the number  $M$  of successive additive approximations.

- 1 Fit initial approximation  $f_0(x)$  (might be taken  $f_0(x) \equiv 0$ )
- 2 For each step  $m = 1, 2, \dots, M$ :

# Gradient boosting

**Input:** training dataset  $(x_n, y_n)$ ,  $n = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$  and the number  $M$  of successive additive approximations.

- ❶ Fit initial approximation  $f_0(x)$  (might be taken  $f_0(x) \equiv 0$ )
- ❷ For each step  $m = 1, 2, \dots, M$ :
  - ❶ calculate gradients:  $g_n = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \Big|_{r=F_{m-1}(x_n)}$

# Gradient boosting

**Input:** training dataset  $(x_n, y_n)$ ,  $n = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$  and the number  $M$  of successive additive approximations.

- ❶ Fit initial approximation  $f_0(x)$  (might be taken  $f_0(x) \equiv 0$ )
- ❷ For each step  $m = 1, 2, \dots, M$ :
  - ❶ calculate gradients:  $g_n = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \Big|_{r=F_{m-1}(x_n)}$
  - ❷ fit  $f_m(\cdot)$  to  $\{(x_n, z_n)\}_{n=1}^N$ , for example by solving

$$\sum_{n=1}^N (f_m(x_n) - g_n)^2 \rightarrow \min_{f_m}$$

# Gradient boosting

**Input:** training dataset  $(x_n, y_n)$ ,  $n = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$  and the number  $M$  of successive additive approximations.

- ❶ Fit initial approximation  $f_0(x)$  (might be taken  $f_0(x) \equiv 0$ )
- ❷ For each step  $m = 1, 2, \dots, M$ :
  - ❶ calculate gradients:  $g_n = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \big|_{r=F_{m-1}(x_n)}$
  - ❷ fit  $f_m(\cdot)$  to  $\{(x_n, z_n)\}_{n=1}^N$ , for example by solving

$$\sum_{n=1}^N (f_m(x_n) - g_n)^2 \rightarrow \min_{f_m}$$

- ❸ solve univariate optimization problem:

$$c_m = \arg \min_{c>0} \sum_{n=1}^N \mathcal{L}(F_{m-1}(x_n) - cf_m(x_n), y_n)$$



# Gradient boosting

**Input:** training dataset  $(x_n, y_n)$ ,  $n = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$  and the number  $M$  of successive additive approximations.

- ① Fit initial approximation  $f_0(x)$  (might be taken  $f_0(x) \equiv 0$ )
- ② For each step  $m = 1, 2, \dots, M$ :
  - ① calculate gradients:  $g_n = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \Big|_{r=F_{m-1}(x_n)}$
  - ② fit  $f_m(\cdot)$  to  $\{(x_n, z_n)\}_{n=1}^N$ , for example by solving

$$\sum_{n=1}^N (f_m(x_n) - g_n)^2 \rightarrow \min_{f_m}$$

- ③ solve univariate optimization problem:

$$c_m = \arg \min_{c > 0} \sum_{n=1}^N \mathcal{L}(F_{m-1}(x_n) - cf_m(x_n), y_n)$$

- ④ set  $F_m(x) = F_{m-1}(x) - c_m f_m(x)$

# Gradient boosting

**Input:** training dataset  $(x_n, y_n)$ ,  $n = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$  and the number  $M$  of successive additive approximations.

- ① Fit initial approximation  $f_0(x)$  (might be taken  $f_0(x) \equiv 0$ )
- ② For each step  $m = 1, 2, \dots, M$ :
  - ① calculate gradients:  $g_n = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \Big|_{r=F_{m-1}(x_n)}$
  - ② fit  $f_m(\cdot)$  to  $\{(x_n, z_n)\}_{n=1}^N$ , for example by solving

$$\sum_{n=1}^N (f_m(x_n) - g_n)^2 \rightarrow \min_{f_m}$$

- ③ solve univariate optimization problem:

$$c_m = \arg \min_{c>0} \sum_{n=1}^N \mathcal{L}(F_{m-1}(x_n) - cf_m(x_n), y_n)$$

- ④ set  $F_m(x) = F_{m-1}(x) - c_m f_m(x)$

**Output:** approximation function  $F_M(x) = f_0(x) - \sum_{m=1}^M c_m f_m(x)$

# Gradient boosting: examples

In gradient boosting

$$\sum_{n=1}^N \left( f_m(x_n) - \frac{\partial \mathcal{L}(r, y)}{\partial r} \Big|_{r=F_{m-1}(x_n)} \right)^2 \rightarrow \min_{f_m}$$

Consider specific cases:

- $\mathcal{L} = \frac{1}{2} (r - y)^2$
- $\mathcal{L} = [-ry]_+$

# Table of Contents

1 Adaboost

2 Gradient boosting

3 Extensions

## Gradient boosting of trees

**Input:** training dataset  $(x_n, y_n)$ ,  $n = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$  and the number  $M$  of successive additive approximations.

- 1 Fit initial approximation *with constant*:

$$f_0(x) = \arg \min_{\gamma} \sum_{n=1}^N \mathcal{L}(\gamma, y_n)$$

## Gradient boosting of trees

**Input:** training dataset  $(x_n, y_n)$ ,  $n = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$  and the number  $M$  of successive additive approximations.

- 1 Fit initial approximation *with constant*:

$$f_0(x) = \arg \min_{\gamma} \sum_{n=1}^N \mathcal{L}(\gamma, y_n)$$

- 2 For each step  $m = 1, 2, \dots, M$ :

## Gradient boosting of trees

**Input:** training dataset  $(x_n, y_n)$ ,  $n = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$  and the number  $M$  of successive additive approximations.

- 1 Fit initial approximation *with constant*:  
$$f_0(x) = \arg \min_{\gamma} \sum_{n=1}^N \mathcal{L}(\gamma, y_n)$$
- 2 For each step  $m = 1, 2, \dots, M$ :
  - 1 calculate gradients  $g_n = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \Big|_{r=F_{m-1}(x_n)}$

## Gradient boosting of trees

**Input:** training dataset  $(x_n, y_n)$ ,  $n = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$  and the number  $M$  of successive additive approximations.

- 1 Fit initial approximation *with constant*:

$$f_0(x) = \arg \min_{\gamma} \sum_{n=1}^N \mathcal{L}(\gamma, y_n)$$

- 2 For each step  $m = 1, 2, \dots, M$ :

- 1 calculate gradients  $g_n = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \Big|_{r=F_{m-1}(x_n)}$
- 2 fit regression tree  $f_m(\cdot)$  on  $\{(x_n, z_n)\}_{n=1}^N$  with some loss function, get leaf regions  $\{R_j^m\}_{j=1}^{J_m}$ .



# Gradient boosting of trees

**Input:** training dataset  $(x_n, y_n)$ ,  $n = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$  and the number  $M$  of successive additive approximations.

- ① Fit initial approximation *with constant*:

$$f_0(x) = \arg \min_{\gamma} \sum_{n=1}^N \mathcal{L}(\gamma, y_n)$$

- ② For each step  $m = 1, 2, \dots, M$ :

- ① calculate gradients  $g_n = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \Big|_{r=F_{m-1}(x_n)}$
- ② fit regression tree  $f_m(\cdot)$  on  $\{(x_n, z_n)\}_{n=1}^N$  with some loss function, get leaf regions  $\{R_j^m\}_{j=1}^{J_m}$ .
- ③ for each terminal region  $R_j^m$ ,  $j = 1, 2, \dots, J_m$  solve univariate optimization problem:

$$\gamma_j^m = \arg \min_{\gamma} \sum_{x_n \in R_j^m} \mathcal{L}(F_{m-1}(x_n) - \gamma, y_n)$$

# Gradient boosting of trees

**Input:** training dataset  $(x_n, y_n)$ ,  $n = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$  and the number  $M$  of successive additive approximations.

- ① Fit initial approximation *with constant*:

$$f_0(x) = \arg \min_{\gamma} \sum_{n=1}^N \mathcal{L}(\gamma, y_n)$$

- ② For each step  $m = 1, 2, \dots, M$ :

- ① calculate gradients  $g_n = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \Big|_{r=F_{m-1}(x_n)}$
- ② fit regression tree  $f_m(\cdot)$  on  $\{(x_n, z_n)\}_{n=1}^N$  with some loss function, get leaf regions  $\{R_j^m\}_{j=1}^{J_m}$ .
- ③ for each terminal region  $R_j^m$ ,  $j = 1, 2, \dots, J_m$  solve univariate optimization problem:

$$\gamma_j^m = \arg \min_{\gamma} \sum_{x_n \in R_j^m} \mathcal{L}(F_{m-1}(x_n) - \gamma, y_n)$$

- ④ update  $F_m(x) = F_{m-1}(x) - \sum_{j=1}^{J_m} \gamma_j^m \mathbb{I}[x \in R_j^m]$

# Gradient boosting of trees

**Input:** training dataset  $(x_n, y_n)$ ,  $n = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$  and the number  $M$  of successive additive approximations.

- ① Fit initial approximation *with constant*:

$$f_0(x) = \arg \min_{\gamma} \sum_{n=1}^N \mathcal{L}(\gamma, y_n)$$

- ② For each step  $m = 1, 2, \dots, M$ :

- ① calculate gradients  $g_n = \frac{\partial \mathcal{L}(r, y_n)}{\partial r} \Big|_{r=F_{m-1}(x_n)}$
- ② fit regression tree  $f_m(\cdot)$  on  $\{(x_n, z_n)\}_{n=1}^N$  with some loss function, get leaf regions  $\{R_j^m\}_{j=1}^{J_m}$ .
- ③ for each terminal region  $R_j^m$ ,  $j = 1, 2, \dots, J_m$  solve univariate optimization problem:

$$\gamma_j^m = \arg \min_{\gamma} \sum_{x_n \in R_j^m} \mathcal{L}(F_{m-1}(x_n) - \gamma, y_n)$$

- ④ update  $F_m(x) = F_{m-1}(x) - \sum_{j=1}^{J_m} \gamma_j^m \mathbb{I}[x \in R_j^m]$

**Output:** approximation function  $F_M(x)$

## Modification of boosting for trees

- Compared to first method of gradient boosting, boosting of regression trees finds additive coefficients individually for each terminal region  $R_j^m$ , not globally for the whole classifier  $f_m(x)$ .
- This is done to increase accuracy: forward stagewise algorithm cannot be applied to find  $R_j^m$ , but it can be applied to find  $\gamma_j^m$ , because second task is solvable for arbitrary  $L$ .
- Max depth  $K$ : interaction between  $K$  features
- Max leaves  $K$ : interaction between no more than  $\leq K - 1$  features
  - usually  $2 \leq K \leq 8$
- $M$  controls underfitting-overfitting trade-off and selected using validation set

# Shrinkage & subsampling

- Shrinkage of general GB, step (d):

$$F_m(x) = F_{m-1}(x) - \alpha c_m f_m(x)$$

- Comments:

- $\alpha \in (0, 1]$
- $\alpha \downarrow \implies M \uparrow (\alpha M \approx \text{const})$

- Subsampling

- increases speed of fitting
- may increase accuracy (diversity of base learners $\uparrow$ )

## Quadratic loss function approximation

$$\begin{aligned}
 \text{Define } g(x) &= \left. \frac{\partial \mathcal{L}(r, y)}{\partial r} \right|_{r=F(x)}, \quad h(x) = \left. \frac{\partial^2 \mathcal{L}(r, y)}{\partial r^2} \right|_{r=F(x)} \\
 \mathcal{L}(F(x) + f(x), y) &\approx \\
 \mathcal{L}(F(x), y) + g(x)f(x) + \frac{1}{2}h(x)(f(x))^2 &= \\
 \frac{1}{2}h(x) \left( f(x) + \frac{g(x)}{h(x)} \right)^2 + \text{const}(f(x))
 \end{aligned}$$

So  $f(x)$  should be fitted to  $-g(x)/h(x)$  with weight  $h(x)$ .

- $h(x) \geq 0$  around local minimum.

Case  $y \in \{1, 2, \dots, C\}$

One-vs-all, one-vs-one, error-correcting-codes.

## Case $y \in \{1, 2, \dots, C\}$

One-vs-all, one-vs-one, error-correcting-codes.

Alternatively can optimize  $\mathcal{L}(F(x), y)$  for  $F(x) \in \mathbb{R}^C$

- $F(x) = \{p(y = c|x)\}_{c=1}^C$ ,  $y$  - one-hot encoded true class
- $S(F(x), y) = F(x)^T y = p(y = \text{correct class}|x)$  - score on  $(x, y)$
- $g_n = -\frac{\partial \mathcal{S}(r, y)}{\partial r} \Big|_{r=F_{m-1}(x_n)} \in \mathbb{R}^C$
- $\sum_{n=1}^N (f_m(x_n) - g_n)^2 \rightarrow \min_{f_m}$  yields vector  $C$ -dim. regression.
- may use quadratic approximation
  - for efficient inverting of  $\left( \frac{\partial^2}{\partial r^2} \mathcal{L}(r, y) \Big|_{r=F(x)} \right)$  may use diagonal approximation.



# xgBoost

- One of the most popular algorithms on kaggle.
- Uses decision trees as base learners:
  - $f_m \in \{f(x) = w_{q(x)}\}$ ,
  - $T$  total number of leaves.
  - $q(x)$  maps  $x \in \mathbb{R}^D$  to leaf number
  - $w \in \mathbb{R}^T$  predictions for leaves.

## xgBoost

- Loss - 2nd order approximation with **with regularization**:

$$\begin{aligned}\mathcal{L}(f_m) &= \sum_{n=1}^N \mathcal{L}(F^{(m-1)}(x_n), y_n) \\ &\approx \sum_{n=1}^N \left[ \mathcal{L}(F^{(m-1)}(x_n), y_n) + g_n f_m(x_n) + \frac{1}{2} h_n f_m^2(x_n) \right] \\ &\quad + \gamma T + \frac{1}{2} \lambda \sum_{t=1}^T w_t^2\end{aligned}$$

- Tree impurity function matches original loss  $\mathcal{L}(\cdot, \cdot)$ .
- Efficiency optimization:
  - feature values may be discretized for speed
  - parallelization over multiple CPU cores and with GPU

# Types of boosting

- Loss function  $\mathcal{L}$ :
  - $\mathcal{L}(|f(x) - y|)$  - regression
  - $F(y \cdot \text{score}(y = +1|x))$  - binary classification
  - $\mathcal{L}(F(x), y)$  for  $F(x), y \in \mathbb{R}^C$  - multiclass classification
- Optimization
  - analytical (Adaboost)
  - gradient based
  - based on quadratic approximation
- Base learners
  - continuous
  - discrete
- Classification
  - binary
  - multiclass
- Extensions: shrinkage, subsampling